

# Adversarial Search II

---

*PROF. LIM KWAN HUI*

50.021 Artificial Intelligence

*The following notes are compiled from various sources such as textbooks, lecture materials, Web resources and are shared for academic purposes only, intended for use by students registered for a specific course. In the interest of brevity, every source is not cited. The compiler of these notes gratefully acknowledges all such sources.*



# Recap: Representing a Game as a Search Problem


---

- We can formally define a strategic two-player game by:
  - Initial State
  - Actions
  - Terminal Test (Win / Lose / Draw)
  - Utility Function (numerical reward for the outcome)
    - Chess: +1, 0, -1
    - Poker: Cash won or lose
- In a zero-sum game with two players
  - each player's utility for a state are equal and opposite



# Recap: Minimax Algorithm

- E.g., 2-ply game



# Recap: Properties of Minimax

---

- Completeness: Yes, if tree is finite
- Optimality: Yes, against an optimal opponent.
- Time complexity:  $O(b^m)$
- Space complexity:  $O(bm)$  (depth-first exploration)
- What issues might there be in terms of time complexity?



# Recap: Properties of Minimax

---

- Completeness: Yes, if tree is finite
- Optimality: Yes, against an optimal opponent.
- Time complexity:  $O(b^m)$
- Space complexity:  $O(bm)$  (depth-first exploration)
  
- What issues might there be in terms of time complexity?
  - For chess,  $b \approx 35$ ,  $m \approx 100$  for “reasonable” games
  - $\Rightarrow$  Exact solution completely infeasible



# Resource Limits

---

- Suppose we have 100 seconds, able to explore  $10^4$  nodes/second
  - $\Rightarrow 10^6$  nodes per move
- Standard Approach:
  - Cutoff test
    - e.g., depth limit
  - Evaluation function
    - = estimated desirability of position



# Resource Limits


---

- Suppose we have 100 seconds, able to explore  $10^4$  nodes/second
  - $\Rightarrow 10^6$  nodes per move
- Standard Approach:
  - Cutoff test
    - e.g., depth limit
  - Evaluation function
    - = estimated desirability of position
- What type of evaluation function can you think of?
  - E.g., for a game of chess




# Evaluation Functions

- For chess, typically linear weighted sum of features
  - $\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
  - e.g.,  $w_1 = 9$  with  $f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$



**Black to move**

White slightly better



**White to move**


Black winning





# Digression: Exact values don't matter

- Behaviour is preserved under any *monotonic* transformation of Eval
- Only the order matters:
  - payoff in deterministic games acts as an *ordinal utility* function



# Cutting off search

- MinimaxCutoff is identical to MinimaxValue except
  - Terminal? is replaced by Cutoff?
  - Utility is replaced by Eval

**Operator = Action or Move**

```
function MINIMAX-DECISION(game) returns an operator
  for each op in OPERATORS[game] do
    VALUE[op] ← MINIMAX-VALUE(APPLY(op, game), game)
  end
  return the op with the highest VALUE[op]
```

**Replaced by Cutoff-Test**

```
function MINIMAX-VALUE(state, game) returns a utility value
  if TERMINAL-TEST[game](state) then
    return UTILITY[game](state)
  else if MAX is to move in state then
    return the highest MINIMAX-VALUE of SUCCESSORS(state)
  else
    return the lowest MINIMAX-VALUE of SUCCESSORS(state)
```

**Replaced by Eval score**



# Cutting off search


---

- MinimaxCutoff is identical to MinimaxValue except
  - Terminal? is replaced by Cutoff?
  - Utility is replaced by Eval
- Does it work in practice?
  - $b^m = 10^6$ ,  $b = 35 \Rightarrow m = 4$
  - 4-ply lookahead is a novice chess player!




# $\alpha$ - $\beta$ Pruning Example

- Original Minimax




# $\alpha$ - $\beta$ Pruning Example

- Intuitive example




# $\alpha$ - $\beta$ Pruning Example

- Intuitive example




# $\alpha$ - $\beta$ Pruning Example

- Intuitive example



# $\alpha$ - $\beta$ Pruning Example


- Intuitive example






# $\alpha$ - $\beta$ Pruning Example

- Intuitive example



# $\alpha$ - $\beta$ Pruning Example

- Intuitive example



# Properties of $\alpha$ - $\beta$ Pruning

---

- Pruning does not affect final result
- Good move ordering improves effectiveness of pruning
- With “perfect ordering”, time complexity is  $O(b^{m/2})$ 
  - $\Rightarrow$  doubles depth of search
  - $\Rightarrow$  can easily reach depth 8 and play good chess
- A simple example of the value of reasoning about which computations are relevant (a form of meta-reasoning)



# $\alpha - \beta$ Algorithm

**function** MAX-VALUE(*state*, *game*,  $\alpha$ ,  $\beta$ ) **returns** the minimax value of *state*

**inputs:** *state*, current state in game

*game*, game description

$\alpha$ , the best score for MAX along the path to *state*

$\beta$ , the best score for MIN along the path to *state*

**if** CUTOFF-TEST(*state*) **then return** EVAL(*state*)

**for each** *s* **in** SUCCESSORS(*state*) **do**

$\alpha \leftarrow \text{MAX}(\alpha, \text{MIN-VALUE}(s, \text{game}, \alpha, \beta))$

**if**  $\alpha \geq \beta$  **then return**  $\beta$

**end**

**return**  $\alpha$

---

**function** MIN-VALUE(*state*, *game*,  $\alpha$ ,  $\beta$ ) **returns** the minimax value of *state*

**if** CUTOFF-TEST(*state*) **then return** EVAL(*state*)

**for each** *s* **in** SUCCESSORS(*state*) **do**

$\beta \leftarrow \text{MIN}(\beta, \text{MAX-VALUE}(s, \text{game}, \alpha, \beta))$

**if**  $\beta \leq \alpha$  **then return**  $\alpha$

**end**


**return**  $\beta$



# $\alpha$ - $\beta$ Pruning Example

MAX

MIN



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$




# $\alpha$ - $\beta$ Pruning Example

*Explore  $a_1$  first*

MAX


MIN



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$




# $\alpha$ - $\beta$ Pruning Example



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$



# $\alpha$ - $\beta$ Pruning Example




- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$





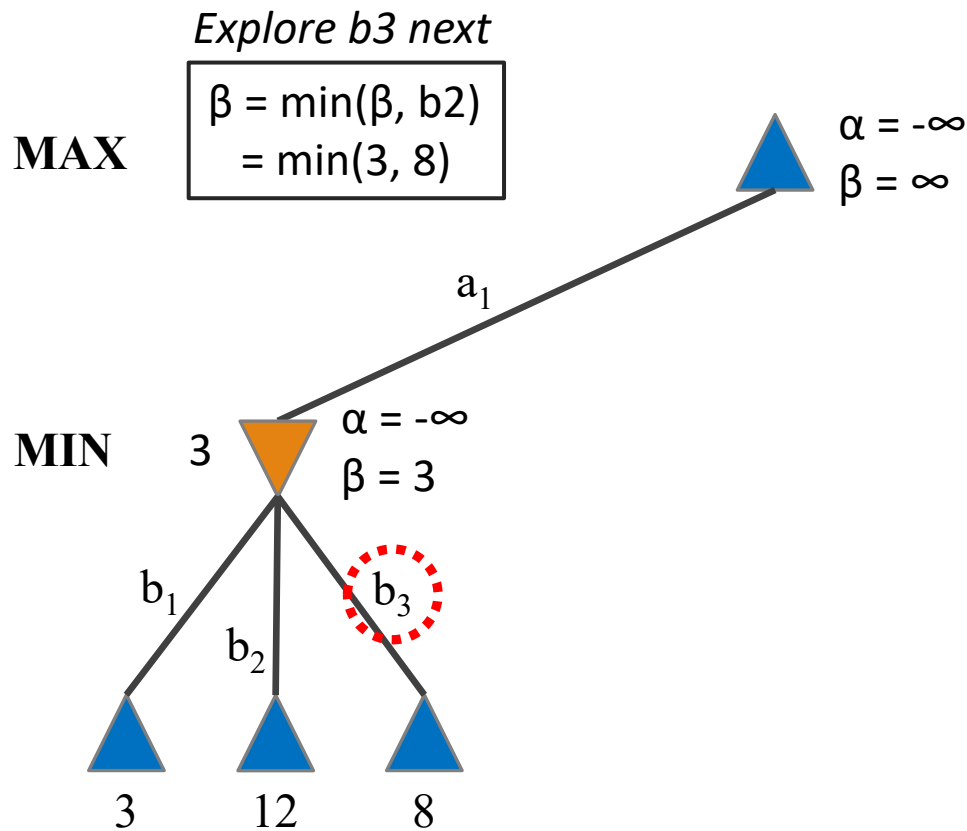
# $\alpha$ - $\beta$ Pruning Example



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$




# $\alpha$ - $\beta$ Pruning Example



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$




# $\alpha$ - $\beta$ Pruning Example



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$




# $\alpha$ - $\beta$ Pruning Example



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$



# $\alpha$ - $\beta$ Pruning Example




- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$



# $\alpha$ - $\beta$ Pruning Example

After exploring  $C_1$ ,  
 **$\alpha \geq \beta$**  so we prune  
the rest, i.e.,  $c_2, c_3$



- $\alpha$  is the best value (to MAX) found so far off the current path
- $\beta$  is the best value (to MIN) found so far off the current path
- Prune if  $\alpha \geq \beta$




# Non-deterministic Games

- Adversarial search where the actions/transitions are non-deterministic
  - E.g., in backgammon, the dice rolls determine the legal moves
- Simplified example with coin-flipping instead of dice rolling

**MAX**

**CHANCE**


**MIN**



# Non-deterministic Games

- **ExpectiMiniMax**, similar to MiniMax but accounts for chance nodes

**if** state is a chance node **then**  
**return** weighted average of ExpectiMinimax-Value of Successors(state)





# Summary & Objectives

---

- Understand the differences between standard search problems and adversarial search problems
- Understand the workings behind the Minimax algorithm and  $\alpha$  -  $\beta$  pruning
- Able to use Minimax algorithm to solve an adversarial/game search problem
- Able to use  $\alpha$  -  $\beta$  pruning to speed up adversarial/game search



# Reminders

---

- AI Quiz 2
  - Time/Venue: Tue 2 Apr 2024, 3.30pm @ LT5
  - Topics Covered: Weeks 5 to 10
  - Format: MCQs and open-ended questions. Completed within 60min
  - 1 x A4 cheatsheet (double-sided) allowed, calculators allowed. Cheatsheet can be printed or handwritten.
  - You will not be asked to produce codes but may be given partial pseudo codes and asked to do something with it, e.g., complete, explain, correct, etc.

