

Planning I

PROF LIM KWAN HUI

50.021 Artificial Intelligence

The following notes are compiled from various sources such as textbooks, lecture materials, Web resources and are shared for academic purposes only, intended for use by students registered for a specific course. In the interest of brevity, every source is not cited. The compiler of these notes gratefully acknowledges all such sources.



Outline & Objectives

- Able to describe planning problems in a simple formalism - that is finding a sequence of actions.
- Able to formulate a planning problem as a STRIPS instance
- Understand the PDDL planning language
- Develop relaxed versions of planning problems and understand the various heuristics that can be used



Planning and Search

- Planning is the process of computing several steps of a problem-solving procedure before executing any of them
- This problem can be solved by search
- The main difference between search and planning is the representation of states
 - In search, states are represented as a single entity (which may be quite a complex object, but its internal structure is not used by the search algorithm)
 - In planning, states have structured representations (collections of properties) which are used by the planning algorithm



Types of Planners

1. Domain-specific

- Made or tuned for a specific planning domain
- Won't work well (if at all) in other planning domains

2. Domain-independent

- In principle, works in any planning domain
- In practice, need restrictions on what kind of planning domain




3. Configurable

- Domain-independent planning engine
- Input includes info about how to solve problems in some domain




Domain-Specific Planners

- Most successful real-world planning systems work this way
 - Mars exploration, sheet-metal bending, playing bridge, etc.
- Often use problem-specific techniques that are difficult to generalize to other planning domains



Domain-Independent Planners

- In principle, works in any planning domain
- No domain-specific knowledge except the description of the system
- In practice,
 - Not feasible to make domain-independent planners work well in all possible planning domains
- Make simplifying assumptions to restrict the set of domains
 - *Classical planning*
 - Historical focus of most research on automated planning



Configurable Planners

- In a fixed planning domain, a domain-independent planner usually won't work as well as a domain-specific planner made for that domain
 - A domain-specific planner may be able to go directly toward a solution in situations where a domain-independent planner would explore many alternative paths
- But we don't want to write a whole new planner for every domain
- Configurable planners
 - Domain-independent planning engine
 - Input includes info about how to solve problems in the domain
- Generally this means one can write a planning engine with fewer restrictions than domain-independent planners



Classical Planning

- Classical planning requires certain assumptions
 - Offline generation of action sequences for an environment that is fully observable, deterministic, finite, static and discrete
- Reduces to the following problem:
 - Given a planning problem $\mathcal{P} = (A, s_0, S_g)$
 - Find a sequence of actions (a_1, a_2, \dots, a_n) that produces a sequence of state transitions (s_1, s_2, \dots, s_n) such that s_n is in S_g .
- This is just path-searching in a graph
 - Nodes = states
 - Edges = actions



STRIPS

- STRIPS = Stanford Research Institute Problem Solver (1971)
 - Originally a planner software, today mostly used to name a formal language to describe planning problems.
 - (Logic-based) Language expressive enough to describe a wide variety of problems, but restrictive enough to allow efficient algorithms to operate over it
- Planning is about finding a sequence of actions to achieve a goal. how to describe such things?




STRIPS Instance

- An initial state;
- The specification of the goal states – situations which the planner is trying to reach
- A set of actions. For each action, the following are included:
 - preconditions: true and false facts that must hold so that an action can be performed
 - postcondition: facts that change when an action is performed.




Example: Robot World

- A simple world with a robot, a ball and two rooms connected by a door
 - How would you model it?




Example: Robot World

- Propositional variables or facts
 - True/false variables that model some aspect of the world
 - $\text{inB}(\text{ball})$: Is the ball in room B? False
 - $\text{inA}(\text{robot})$: Is the robot in room A? True




Example: Robot World

- There are various objects (robot, ball) in this world
 - How do we model actions on these objects, e.g., kick ball to other room?
 - What effect does it have on the variables representing these objects?




Example: Robot World

- There are various objects (robot, ball) in this world
 - Actions may change the value of these variables
 - Actions may have various preconditions of variable values




Example: Robot World

- For modelling the action of the robot kicking the ball to room B, $\text{kickball}(a,b)$, what do we need check for first?




Example: Robot World

- For modelling the action of the robot kicking the ball to room B, $\text{kickball}(a,b)$, what do we need check for first?
 - $\text{inA}(\text{robot})$, $\text{inA}(\text{ball})$, $\text{dooropen}(A,B)$




Example: Robot World

- For modelling the action of the robot kicking the ball to room B, `kickball(a,b)`, what changes are there to the variables?



Example: Robot World

- For modelling the action of the robot kicking the ball to room B, $\text{kickball}(a,b)$, what changes are there to the variables?
 - $\text{not inA}(\text{ball})$, $\text{inB}(\text{ball})$
 - What about the robot? Still $\text{inA}(\text{robot})$



STRIPS Instance

- An initial state;
- The specification of the goal states – situations which the planner is trying to reach
- A set of actions. For each action, the following are included:
 - preconditions: true and false facts that must hold so that an action can be performed
 - postcondition: facts that change when an action is performed.



STRIPS Instance

- STRIPS is formally defined as a 4-tuple (P, O, I, G)
- P - a set of propositional variables – the facts that describe the state of the world
- O - a set of operators (i.e., actions). Each operator itself is a 3-tuple (pre_a, add_a, del_a) of
 - pre_a - facts that must be true before the action can be performed
 - add_a - facts that will change to true when/after the action can be performed
 - del_a - facts that will change to false when/after the action can be performed
 - requirement: $add_a \cap del_a = \emptyset$
- I - the initial state of the world, true/false assignments to variables from P
- G - the goal state of the world




STRIPS: Robot World

- STRIPS is formally defined as a 4-tuple (P, O, I, G)
- P - a set of propositional variables. E.g., $\text{inA}(x)$, $\text{inB}(x)$, $\text{dooropen}(x,y)$
- O - a set of operators (i.e., actions). E.g., $\text{kickball}(a,b)$
 - pre_a - $\text{inA}(\text{robot})$, $\text{inA}(\text{ball})$, $\text{dooropen}(A,B)$
 - add_a - $\text{inB}(\text{ball})$
 - del_a - $\text{inA}(\text{ball})$
- I - the initial state of the world. E.g., $\text{inA}(\text{robot})$, $\text{inA}(\text{ball})$, $\text{dooropen}(A,B)$
- G - the goal state of the world. E.g., $\text{inB}(\text{ball})$



State Representation

- World states are represented as sets of facts: conjunction of propositions (conditions)
 - E.g., Robot World: State 1 = { $\text{inA}(\text{robot}) \wedge \text{inA}(\text{ball}) \wedge \text{dooropen}(\text{A},\text{B})$ }



- Closed World Assumption (CWA): Facts not listed in a state are assumed to be false. Under CWA the assumption the agent has full observability and only positive facts need to be stated



State Representation

- The world is represented through a set of features / objects (e.g., planes, people, cities) and each proposition states a fact that attributes “values” to features

Objects	State Propositions	Closed World Assumptions
Robot	inA(robot)	Not inB(robot)
Ball	inA(ball)	Not inB(ball)
Door	doorOpen(a,b)	



State Representation



- Goals states are also represented as sets of facts
 - E.g., the state $\{ \text{inB}(\text{ball}) \}$ can be defined as a goal state
- A goal state is any state that includes all the goal facts
 - The following are valid goals:
 - State 1 = $\{ \text{inA}(\text{robot}) \wedge \text{inB}(\text{ball}) \}$
 - State 2 = $\{ \text{inB}(\text{robot}) \wedge \text{inB}(\text{ball}) \}$
 - State 3 = $\{ \text{dooropen}(a,b) \wedge \text{inB}(\text{robot}) \wedge \text{inB}(\text{ball}) \}$
 - The following are not:
 - State 1 = $\{ \text{inB}(\text{robot}) \wedge \text{dooropen}(a,b) \}$
 - State 2 = $\{ \text{inA}(\text{ball}) \wedge \text{inB}(\text{robot}) \}$





Exercise: Block World

- A simple world that comprises a set of blocks, a table, and a robot claw
- The aim is to stack the blocks on top of each other in alphabetical order
- Some constraints of this world:
 - The robot claw can only carry one block at a time
 - Only one block can be on another block
 - Any number of blocks can be on the table
- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P, I and G first.

Initial State





Goal State




Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P , I and G first.
- **P** - a set of propositional variables.
 - ?

Initial State





Goal State




Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P , I and G first.
- **P** - a set of propositional variables.
 - $\text{ontable}(\text{block1})$: block1 is on the table
 - $\text{on}(\text{block1}, \text{block2})$: block1 is on block2
 - $\text{free}(\text{block1})$: top of block1 is free
 - clawempty : the robot claw is not holding any blocks

Initial State





Goal State




Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P , I and G first.
- **P** - a set of propositional variables.
 - $\text{ontable}(\text{block1})$: block1 is on the table
 - $\text{on}(\text{block1}, \text{block2})$: block1 is on block2
 - $\text{free}(\text{block1})$: top of block1 is free
 - clawempty : the robot claw is not holding any blocks
- **I** - the initial state of the world.
 - ?

Initial State





Goal State




Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P , I and G first.
- **P** - a set of propositional variables.
 - $\text{ontable}(\text{block1})$: block1 is on the table
 - $\text{on}(\text{block1}, \text{block2})$: block1 is on block2
 - $\text{free}(\text{block1})$: top of block1 is free
 - clawempty : the robot claw is not holding any blocks
- **I** - the initial state of the world.
 - $\text{ontable}(c)$, $\text{ontable}(a)$, ... , $\text{free}(c)$, ... , clawempty

Initial State





Goal State




Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P , I and G first.
- **P** - a set of propositional variables.
 - $\text{ontable}(\text{block1})$: block1 is on the table
 - $\text{on}(\text{block1}, \text{block2})$: block1 is on block2
 - $\text{free}(\text{block1})$: top of block1 is free
 - clawempty : the robot claw is not holding any blocks
- **I** - the initial state of the world.
 - $\text{ontable}(c)$, $\text{ontable}(a)$, ... , $\text{free}(c)$, ... , clawempty
- **G** - the goal state of the world.
 - ?

Initial State




Goal State




Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Let's start with P , I and G first.
- **P** - a set of propositional variables.
 - $\text{ontable}(\text{block1})$: block1 is on the table
 - $\text{on}(\text{block1}, \text{block2})$: block1 is on block2
 - $\text{free}(\text{block1})$: top of block1 is free
 - clawempty : the robot claw is not holding any blocks
- **I** - the initial state of the world.
 - $\text{ontable}(c)$, $\text{ontable}(a)$, ... , $\text{free}(c)$, ... , clawempty
- **G** - the goal state of the world.
 - $\text{on}(a,b)$, $\text{on}(b,c)$, $\text{on}(c,d)$, $\text{ontable}(d)$

Initial State




Goal State



Exercise: Block World



- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Next, let's move on to O , specifically `grabandstack()`, which grabs a block and stacks it in the last column.
- **O** - a set of operators (i.e., actions). E.g., for the `grabandstack(block1,block2)` operator:
 - pre_a -
 - add_a -
 - del_a -





Exercise: Block World

- Task: How would you model this world in terms of the STRIPS 4-tuple (P, O, I, G) ?
 - Next, let's move on to O , specifically `grabandstack()`, which grabs a block and stacks it in the last column.
- **O** - a set of operators (i.e., actions). E.g., for the `grabandstack(block1, block2)` operator:
 - pre_a – `free(b1), free(b2), clawempty, ontable(b1)`
 - add_a – `on(b1, b2)`
 - del_a – `free(b2), ontable(b1)`

Initial State





Goal State





Exercise: Block World

- Task: Instead of grabandstack(), we want separate grab(block) and stack(block1, block2) operators.
 - What new propositional variables or facts do we need?
 - How do you define the operators?
- **P** - a set of propositional variables.
 - Existing: ontable(blk1), on(b1,b2), free(b1), clawempty
 - New:
- **O** - grab(block):
 - $pre_a -$
 - $add_a -$
 - $del_a -$
- **O** - stack(block1, block2):
 - $pre_a -$
 - $add_a -$
 - $del_a -$

Initial State





Goal State





Exercise: Block World

- Task: Instead of grabandstack(), we want separate grab(block) and stack(block1, block2) operators.
 - What new propositional variables or facts do we need?
 - How do you define the operators?
- **P** - a set of propositional variables.
 - Existing: ontable(blk1), on(b1,b2), free(b1), clawempty
 - New: holding(block1)
- **O** - grab(block):
 - pre_a – ontable(b), free(b), clawempty
 - add_a – holding(b)
 - del_a – ontable(b), free(b), clawempty
- **O** - stack(block1, block2):
 - pre_a – holding(b1), free(b2)
 - add_a – clawempty, free(b1) on(b1,b2)
 - del_a – holding(b1), free(b2)

Initial State



Goal State



Next

- Look at the PDDL planning language
- Simple coding exercise on using PDDL

