

# Parametric Language Modeling : Transformers

# Parametric Language Modeling




- Can we do better than tf-idf?
- We need parametric language modeling.
  - Set a learnable objective.
  - Use this objective to tune the parameters with gradient descent.
- Question is – what network to use?
  - RNN?
  - Transformer?

# Lecture Plan

1. From recurrence (RNN) to attention-based NLP models
2. Introducing the Transformer model
3. Great results with Transformers

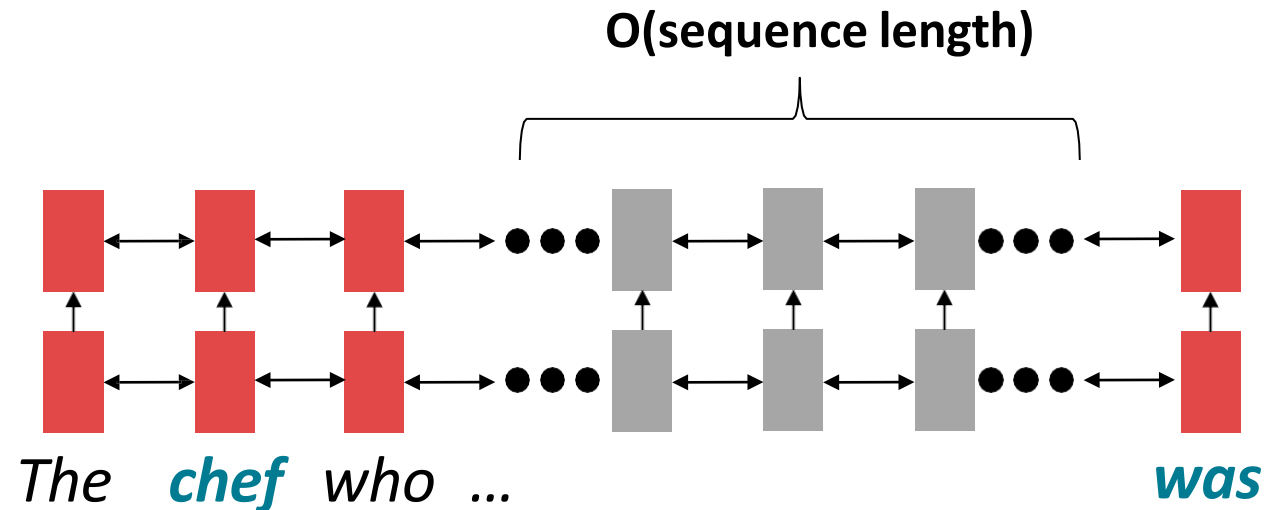

# As of last week: recurrent models for (most) NLP!

- Circa 2016, the de facto strategy in NLP is to **encode** sentences with a bidirectional LSTM: (for example, the source sentence in a translation)
- Define your output (parse, sentence, summary) as a sequence, and use an LSTM to generate it.
- Use attention to allow flexible access to memory




# Issues with recurrent models: Linear interaction distance

- RNNs are unrolled “left-to-right”.
- This encodes linear locality: a useful heuristic!
  - Nearby words often affect each other’s meanings
- **Problem:** RNNs take  $O(\text{sequence length})$  steps for distant word pairs to interact.



# Issues with recurrent models: Linear interaction distance


- **O(sequence length)** steps for distant word pairs to interact means:
  - Hard to learn long-distance dependencies (because of gradient problems!)
  - Linear order of words is “baked in”; we already know linear order isn’t the right way to think about sentences...



Info of *chef* has gone through  $O(\text{sequence length})$  many layers!

# Issues with recurrent models: Lack of parallelizability


- Forward and backward passes have  **$O(\text{sequence length})$**  unparallelizable operations
  - GPUs can perform a bunch of independent computations at once!
  - But future RNN hidden states can't be computed in full before past RNN hidden states have been computed
  - Inhibits training on very large datasets!



Numbers indicate min # of steps before a state can be computed

# If not recurrence, then what? How about attention?

- **Attention** treats each word's representation as a **query** to access and incorporate information from a **set of values**.
- Number of unparallelizable operations does not increase sequence length.
- Maximum interaction distance:  $O(1)$ , since all words interact at every layer!



All words attend to all words in previous layer; most arrows here are omitted



# Self-Attention

- Recall: Attention operates on **queries**, **keys**, and **values**.
  - We have some **queries**  $q_1, q_2, \dots, q_T$ . Each query is  $q_i \in \mathbb{R}^d$
  - We have some **keys**  $k_1, k_2, \dots, k_T$ . Each key is  $k_i \in \mathbb{R}^d$
  - We have some **values**  $v_1, v_2, \dots, v_T$ . Each value is  $v_i \in \mathbb{R}^d$
- In **self-attention**, the queries, keys, and values are drawn from the same source.
  - For example, if the output of the previous layer is  $x_1, \dots, x_T$ , (one vec per word) we could let  $v_i = k_i = q_i = x_i$  (that is, use the same vectors for all of them!)
- The (dot product) self-attention operation is as follows:

The number of queries can differ from the number of keys and values in practice.

$$e_{ij} = q_i^\top k_j$$

Compute **key-query** affinities

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{j'} \exp(e_{ij'})}$$

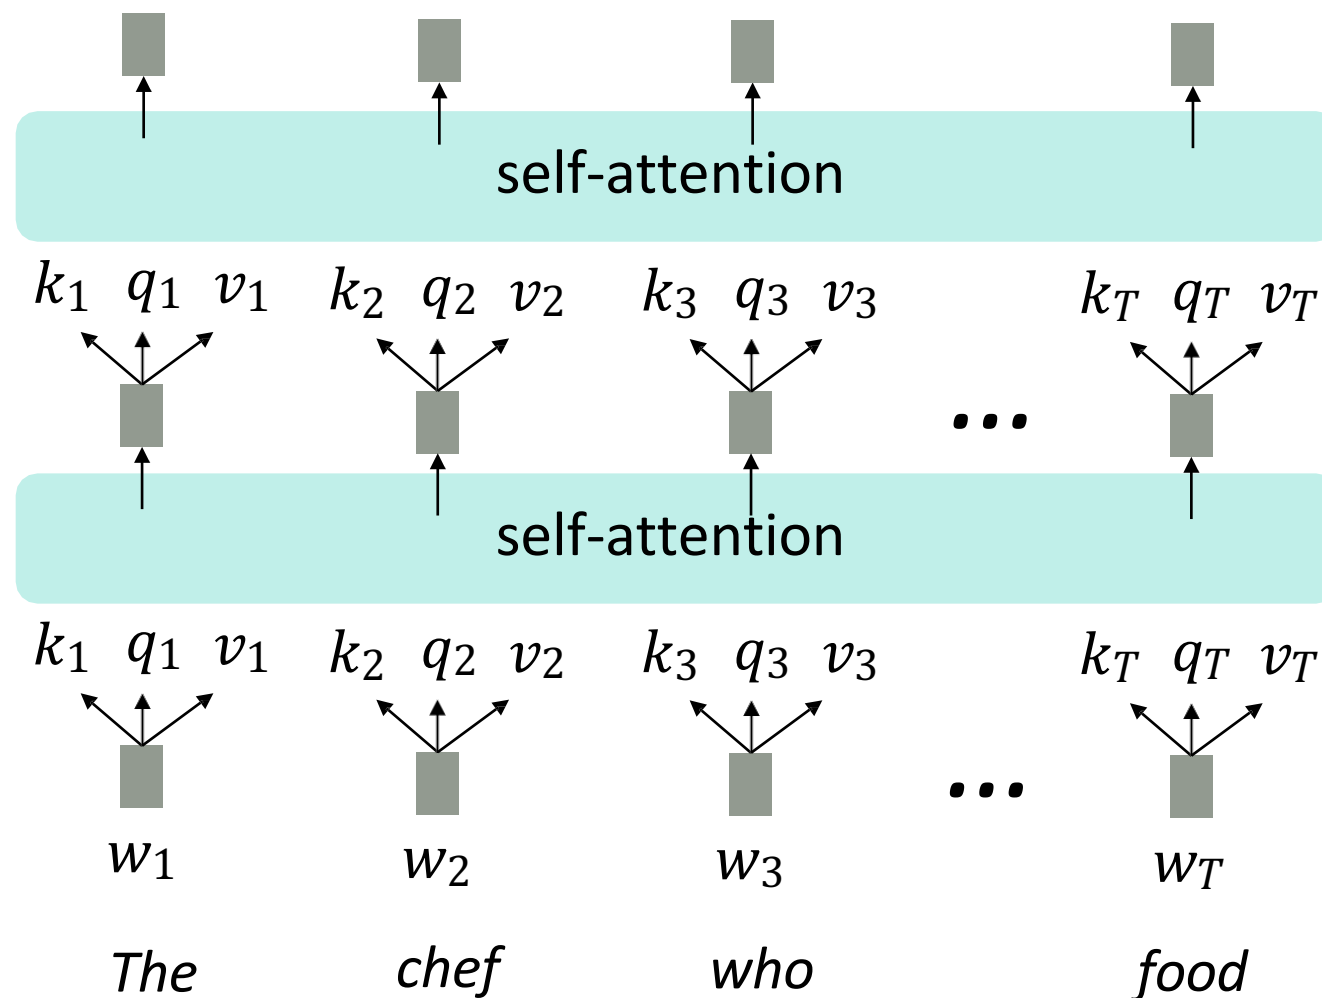
Compute attention weights from affinities (softmax)

$$\text{output}_i = \sum_j \alpha_{ij} v_j$$

Compute outputs as weighted sum of **values**

# Self-attention as an NLP building block

- In the diagram at the right, we have stacked self-attention blocks, like we might stack LSTM layers.
- Can self-attention be a drop-in replacement for recurrence?
- No. It has a few issues, which we'll go through.
- First, self-attention is an operation on **sets**. It has no inherent notion of order.



Self-attention doesn't know the order of its inputs.

# Barriers and solutions for Self-Attention as a building block

## Barriers

- Doesn't have an inherent notion of order!



## Solutions

# Fixing the first self-attention problem: **sequence order**

- Since self-attention doesn't build in order information, we need to encode the order of the sentence in our keys, queries, and values.
- Consider representing each **sequence index** as a **vector**

$p_i \in \mathbb{R}^d$ , for  $i \in \{1, 2, \dots, T\}$  are position vectors

- Don't worry about what the  $p_i$  are made of yet!
- Easy to incorporate this info into our self-attention block: just add the  $p_i$  to our inputs!
- Let  $\tilde{v}_i, \tilde{k}_i, \tilde{q}_i$  be our old values, keys, and queries.

$$v_i = \tilde{v}_i + p_i$$

$$q_i = \tilde{q}_i + p_i$$


$$k_i = \tilde{k}_i + p_i$$

In deep self-attention networks, we do this at the first layer! You could concatenate them as well, but people mostly just add...

# Position representation vectors through sinusoids

- **Sinusoidal position representations:** concatenate sinusoidal functions of varying periods:

$$p_i = \begin{pmatrix} \sin(i/10000^{2*1/d}) \\ \cos(i/10000^{2*1/d}) \\ \vdots \\ \sin(i/10000^{2*\frac{d}{2}/d}) \\ \cos(i/10000^{2*\frac{d}{2}/d}) \end{pmatrix}$$



- Pros:
  - Periodicity indicates that maybe “absolute position” isn’t as important
  - Maybe can extrapolate to longer sequences as periods restart!
- Cons:
  - Not learnable; also the extrapolation doesn’t really work!

# Position representation vectors learned from scratch

- **Learned absolute position representations:** Let all  $p_i$  be learnable parameters!  
Learn a matrix  $p \in \mathbb{R}^{d \times T}$ , and let each  $p_i$  be a column of that matrix!
- Pros:
  - Flexibility: each position gets to be learned to fit the data
- Cons:
  - Definitely can't extrapolate to indices outside  $1, \dots, T$ .
- Most systems use this!
- Sometimes people try more flexible representations of position:
  - Relative linear position attention [\[Shaw et al., 2018\]](#)
  - Dependency syntax-based position [\[Wang et al., 2019\]](#)

# Barriers and solutions for Self-Attention as a building block

## Barriers

- Doesn't have an inherent notion of order!
- No nonlinearities for deep learning! It's all just weighted averages




## Solutions

- Add position representations to the inputs

# Adding nonlinearities in self-attention

- Note that there are no elementwise nonlinearities in self-attention; stacking more self-attention layers just re-averages **value** vectors
- Easy fix: add a **feed-forward network** to post-process each output vector.

$$\begin{aligned} m_i &= MLP(\text{output}_i) \\ &= W_2 * \text{ReLU}(W_1 \times \text{output}_i + b_1) + b_2 \end{aligned}$$



Intuition: the FF network processes the result of attention



# Barriers and solutions for Self-Attention as a building block

## Barriers

- Doesn't have an inherent notion of order!
- No nonlinearities for deep learning magic! It's all just weighted averages
- Need to ensure we don't "look at the future" when predicting a sequence
  - Like in machine translation
  - Or language modeling



## Solutions

- Add position representations to the inputs
- Easy fix: apply the same feedforward network to each self-attention output.

# Masking the future in self-attention

- To use self-attention in **decoders**, we need to ensure we can't peek at the future.
- At every timestep, we could change the set of **keys and queries** to include only past words. (Inefficient!)
- To enable parallelization, we **mask out attention** to future words by setting attention scores to  $-\infty$ .

$$e_{ij} = \begin{cases} q_i^\top k_j, & j < i \\ -\infty, & j \geq i \end{cases}$$

For encoding these words

[START]

We can look at these  
(not greyed out) words

[START]      The      chef      who

[The matrix of  $e_{ij}$  values]

# Masking the future in self-attention

- To use self-attention in **decoders**, we need to ensure we can't peek at the future.
- At every timestep, we could change the set of **keys and queries** to include only past words. (Inefficient!)
- To enable parallelization, we **mask out attention** to future words by setting attention scores to  $-\infty$ .

$$e_{ij} = \begin{cases} q_i^\top k_j, j < i \\ -\infty, j \geq i \end{cases}$$

For encoding these words

We can look at these (not greyed out) words

|         | [START]   | The       | chef      | who       |
|---------|-----------|-----------|-----------|-----------|
| [START] | $-\infty$ | $-\infty$ | $-\infty$ | $-\infty$ |
| The     |           | $-\infty$ | $-\infty$ | $-\infty$ |
| chef    |           |           | $-\infty$ | $-\infty$ |
| who     |           |           |           | $-\infty$ |

# Barriers and solutions for Self-Attention as a building block

## Barriers

- Doesn't have an inherent notion of order!
- No nonlinearities for deep learning magic! It's all just weighted averages
- Need to ensure we don't "look at the future" when predicting a sequence
  - Like in machine translation
  - Or language modeling



## Solutions

- Add position representations to the inputs
- Easy fix: apply the same feedforward network to each self-attention output.
- Mask out the future by artificially setting attention weights to 0!

# Necessities for a self-attention building block:

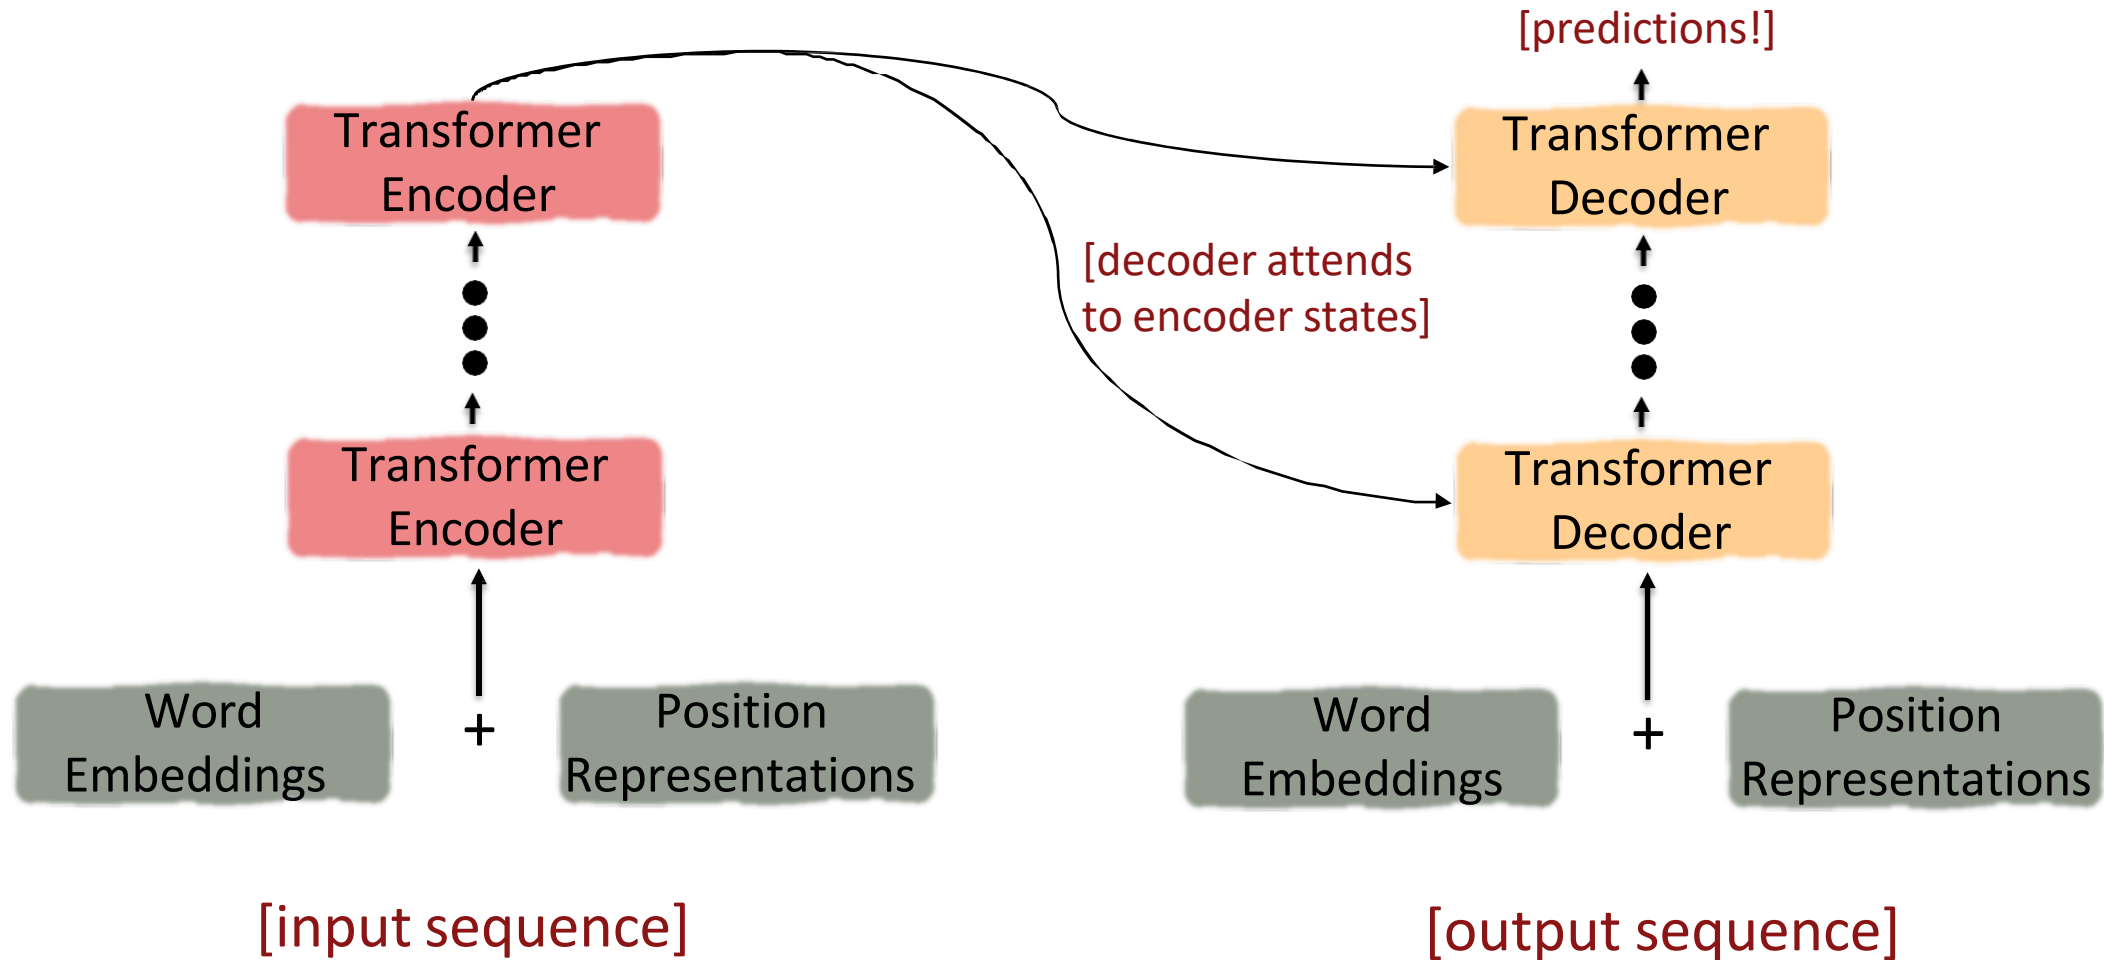
- **Self-attention:**
  - the basis of the method.
- **Position representations:**
  - Specify the sequence order, since self-attention is an unordered function of its inputs.
- **Nonlinearities:**
  - At the output of the self-attention block
  - Frequently implemented as a simple feed-forward network.
- **Masking:**
  - In order to parallelize operations while not looking at the future.
  - Keeps information about the future from “leaking” to the past.
- That’s it! But this is not the **Transformer** model we’ve been hearing about.

# Outline

1. From recurrence (RNN) to attention-based NLP models
2. Introducing the Transformer model
3. Great results with Transformers

# The Transformer Encoder-Decoder [\[Vaswani et al., 2017\]](#)

First, let's look at the Transformer Encoder and Decoder Blocks at a high level



# The Transformer Encoder-Decoder [\[Vaswani et al., 2017\]](#)

Next, let's look at the Transformer Encoder and Decoder Blocks

What's left in a Transformer Encoder Block that we haven't covered?

1. **Key-query-value attention:** How do we get the  $k, q, v$  vectors from a single word embedding?
2. **Multi-headed attention:** Attend to multiple places in a single layer!
3. **Tricks to help with training!**
  1. Residual connections
  2. Layer normalization
  3. Scaling the dot product
  4. These tricks **don't improve** what the model is able to do; they help improve the training process. Both of these types of modeling improvements are very important!



# The Transformer Encoder: Key-Query-Value Attention

- We saw that self-attention is when keys, queries, and values come from the same source. The Transformer does this in a particular way:
  - Let  $x_1, \dots, x_T$  be input vectors to the Transformer encoder;  $x_i \in \mathbb{R}^d$
- Then keys, queries, values are:
  - $k_i = Kx_i$ , where  $K \in \mathbb{R}^{d \times d}$  is the key matrix.
  - $q_i = Qx_i$ , where  $Q \in \mathbb{R}^{d \times d}$  is the query matrix.
  - $v_i = Vx_i$ , where  $V \in \mathbb{R}^{d \times d}$  is the value matrix.
- These matrices allow *different aspects* of the  $x$  vectors to be used/emphasized in each of the three roles.

# The Transformer Encoder: Key-Query-Value Attention

- Let's look at how key-query-value attention is computed, in matrices.
  - Let  $X = [x_1; \dots; x_T] \in \mathbb{R}^{T \times d}$  be the concatenation of input vectors.
  - First, note that  $XK \in \mathbb{R}^{T \times d}$ ,  $XQ \in \mathbb{R}^{T \times d}$ ,  $XV \in \mathbb{R}^{T \times d}$ .
  - The output is defined as  $\text{output} = \text{softmax}(XQ(XK)^T) \times XV$ .

First, take the query-key dot products in one matrix multiplication:  $XQ(XK)^T$

A diagram illustrating the first step of the attention mechanism. It shows a vertical pink box labeled  $XQ$  multiplied by a horizontal pink box labeled  $K^T X^T$ . The result is a larger vertical pink box labeled  $XQK^T X^T$ , followed by  $\in \mathbb{R}^{T \times T}$ . To the right of this box, the text "All pairs of attention scores!" is written in blue.

$$XQ \cdot K^T X^T = XQK^T X^T \in \mathbb{R}^{T \times T}$$

All pairs of attention scores!

Next, softmax, and compute the weighted average with another matrix multiplication.

A diagram illustrating the second step of the attention mechanism. It shows the expression  $\text{softmax} \left( XQK^T X^T \right) \cdot XV = \text{output}$ . The matrix  $XQK^T X^T$  is enclosed in large parentheses, and the word "softmax" is written to its left. The result is a vertical pink box labeled "output", followed by  $\in \mathbb{R}^{T \times d}$ . An arrow points from the  $XQK^T X^T$  box in the previous diagram to the  $XQK^T X^T$  box in this diagram.



$$\text{softmax} \left( XQK^T X^T \right) \cdot XV = \text{output} \in \mathbb{R}^{T \times d}$$

# The Transformer Encoder: **Multi-headed attention**

- What if we want to look in multiple places in the sentence at once?
  - For word  $i$ , self-attention “looks” where  $x_i^\top Q^\top K x_j$  is high, but maybe we want to focus on different  $j$  for different reasons?
- We’ll define **multiple attention “heads”** through multiple Q,K,V matrices
- Let,  $Q_\ell, K_\ell, V_\ell \in \mathbb{R}^{d \times \frac{d}{h}}$ , where  $h$  is the number of attention heads, and  $\ell$  ranges from 1 to  $h$ .
- Each attention head performs attention independently:
  - $\text{output}_\ell = \text{softmax}(X Q_\ell K_\ell^\top X^\top) * X V_\ell$ , where  $\text{output}_\ell \in \mathbb{R}^{d/h}$
- Then the outputs of all the heads are combined!
  - $\text{output} = Y[\text{output}_1; \dots; \text{output}_h]$ , where  $Y \in \mathbb{R}^{d \times d}$
- Each head gets to “look” at different things, and construct value vectors differently.

# The Transformer Encoder: Multi-headed attention


- What if we want to look in multiple places in the sentence at once?
  - For word  $i$ , self-attention “looks” where  $x_i^\top Q^\top K x_j$  is high, but maybe we want to focus on different  $j$  for different reasons?
- We’ll define **multiple attention “heads”** through multiple Q,K,V matrices
- Let,  $Q_\ell, K_\ell, V_\ell \in \mathbb{R}^{d \times \frac{d}{h}}$ , where  $h$  is the number of attention heads, and  $\ell$  ranges from 1 to  $h$ .




Same amount of computation as single-head self-attention!

# The Transformer Encoder: **Residual connections** [[He et al., 2016](#)]


- **Residual connections** are a trick to help models train better.
  - Instead of  $X^{(i)} = \text{Layer}(X^{(i-1)})$  (where  $i$  represents the layer)




- We let  $X^{(i)} = X^{(i-1)} + \text{Layer}(X^{(i-1)})$  (so we only have to learn “the residual” from the previous layer)



- Residual connections are thought to make the loss landscape considerably smoother (thus easier training!)



[no residuals]



[residuals]

[Loss landscape visualization,  
[Li et al., 2018](#), on a ResNet]

# The Transformer Encoder: **Layer normalization** [Ba et al., 2016]

- **Layer normalization** is a trick to help models train faster.
- Idea: cut down on uninformative variation in hidden vector values by normalizing to unit mean and standard deviation **within each layer**.
  - LayerNorm's success may be due to its normalizing gradients [Xu et al., 2019]
- Let  $x \in \mathbb{R}^d$  be an individual (word) vector in the model.
- Let  $\mu = \sum_{j=1}^d x_j$ ; this is the mean;  $\mu \in \mathbb{R}$ .
- Let  $\sigma = \sqrt{\frac{1}{d} \sum_{j=1}^d (x_j - \mu)^2}$ ; this is the standard deviation;  $\sigma \in \mathbb{R}$ .
- Let  $\gamma \in \mathbb{R}^d$  and  $\beta \in \mathbb{R}^d$  be learned “gain” and “bias” parameters. (Can omit!)
- Then layer normalization computes:

Normalize by scalar  
mean and variance

$$\text{output} = \frac{x - \mu}{\sqrt{\sigma} + \epsilon}$$

# The Transformer Encoder: **Layer normalization** [Ba et al., 2016]

- **Layer normalization** is a trick to help models train faster.
- Idea: cut down on uninformative variation in hidden vector values by normalizing to unit mean and standard deviation **within each layer**.
  - LayerNorm's success may be due to its normalizing gradients [Xu et al., 2019]
- Let  $x \in \mathbb{R}^d$  be an individual (word) vector in the model.
- Let  $\mu = \sum_{j=1}^d x_j$ ; this is the mean;  $\mu \in \mathbb{R}$ .
- Let  $\sigma = \sqrt{\frac{1}{d} \sum_{j=1}^d (x_j - \mu)^2}$ ; this is the standard deviation;  $\sigma \in \mathbb{R}$ .
- Let  $\gamma \in \mathbb{R}^d$  and  $\beta \in \mathbb{R}^d$  be learned “gain” and “bias” parameters. (Can omit!)
- Then layer normalization computes:

$$\text{output} = \frac{x - \mu}{\sqrt{\sigma} + \epsilon} * \gamma + \beta$$

Normalize by scalar mean and variance

Modulate by learned elementwise gain and bias

# The Transformer Encoder: **Scaled Dot Product** [Vaswani et al., 2017]

- **“Scaled Dot Product”** attention is a final variation to aid in Transformer training.
- When dimensionality  $d$  becomes large, dot products between vectors tend to become large.
  - Because of this, inputs to the softmax function can be large, making the gradients small.

- Instead of the self-attention function we’ve seen:

$$\text{output}_\ell = \text{softmax}(XQ_\ell K_\ell^\top X^\top) * XV_\ell$$

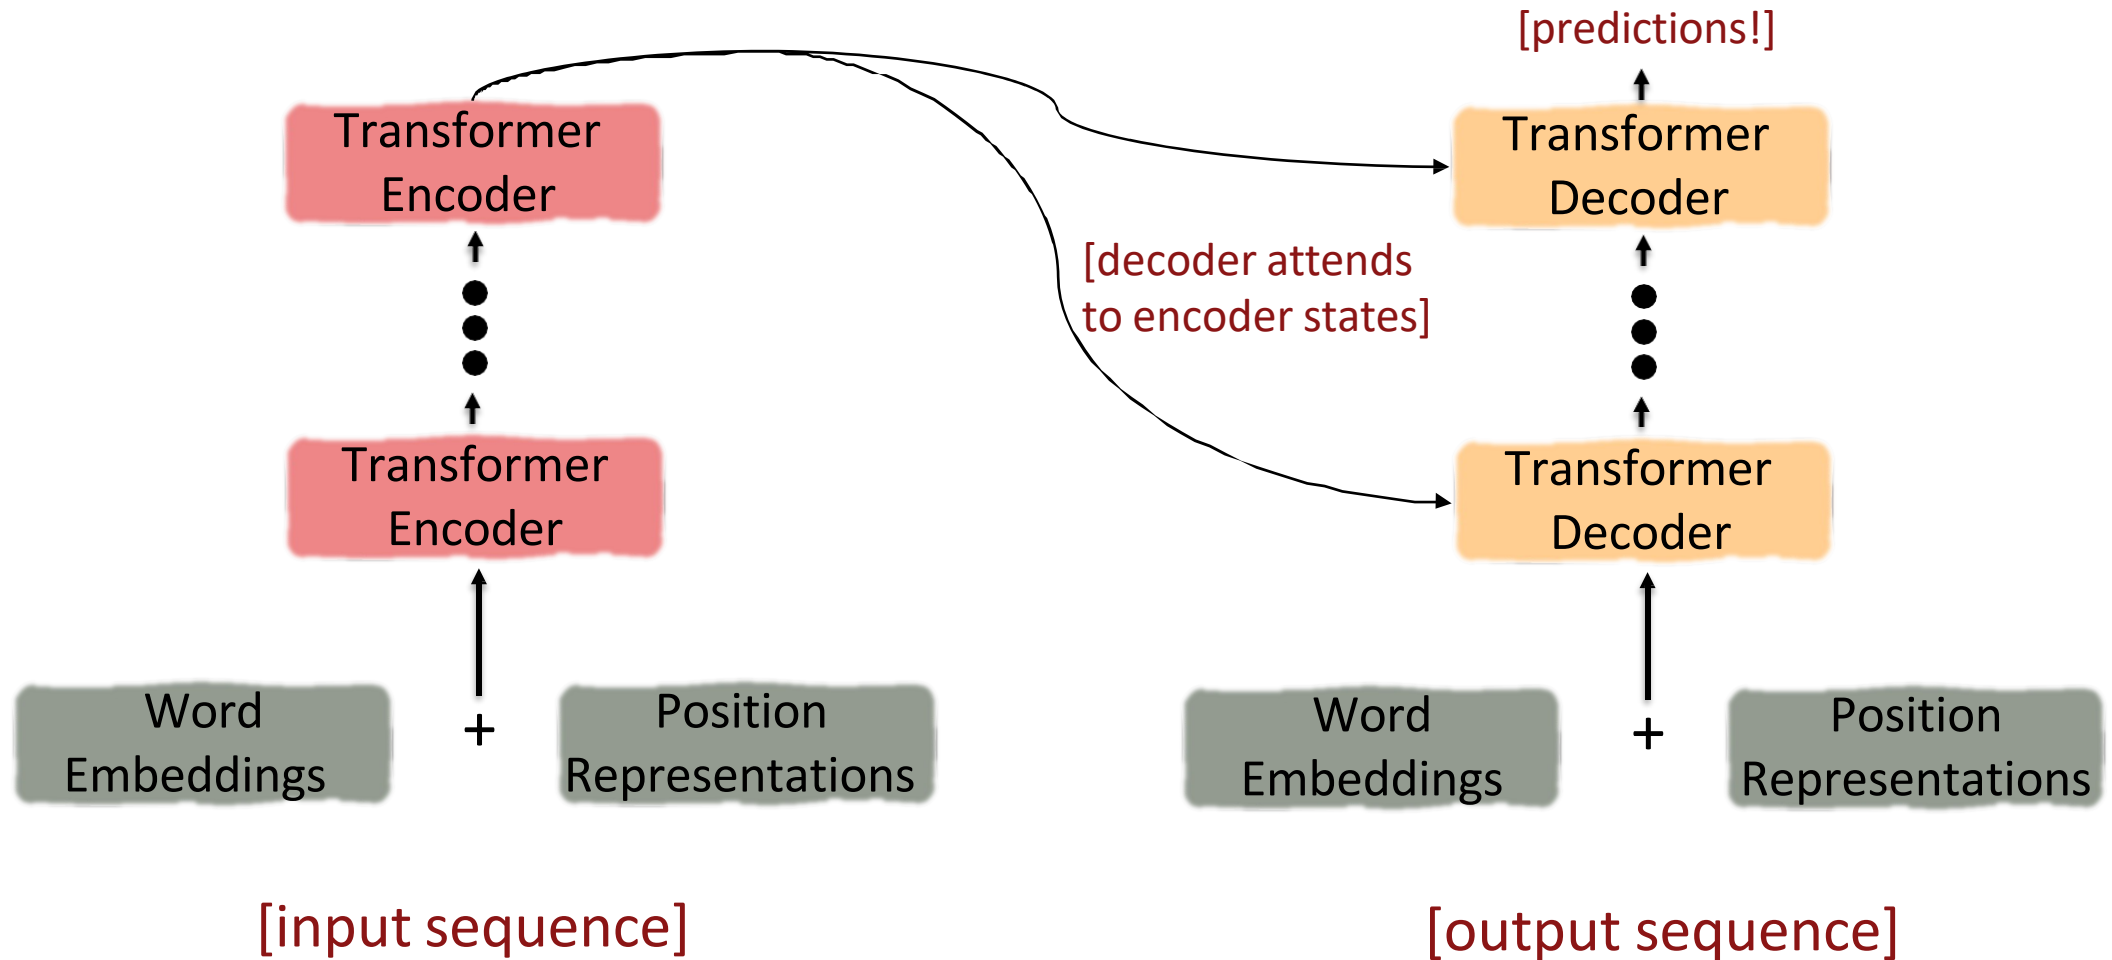
- We divide the attention scores by  $\sqrt{d/h}$ , to stop the scores from becoming large just as a function of  $d/h$  (The dimensionality divided by the number of heads.)

$$\text{output}_\ell = \text{softmax}\left(\frac{XQ_\ell K_\ell^\top X^\top}{\sqrt{d/h}}\right) * XV_\ell$$



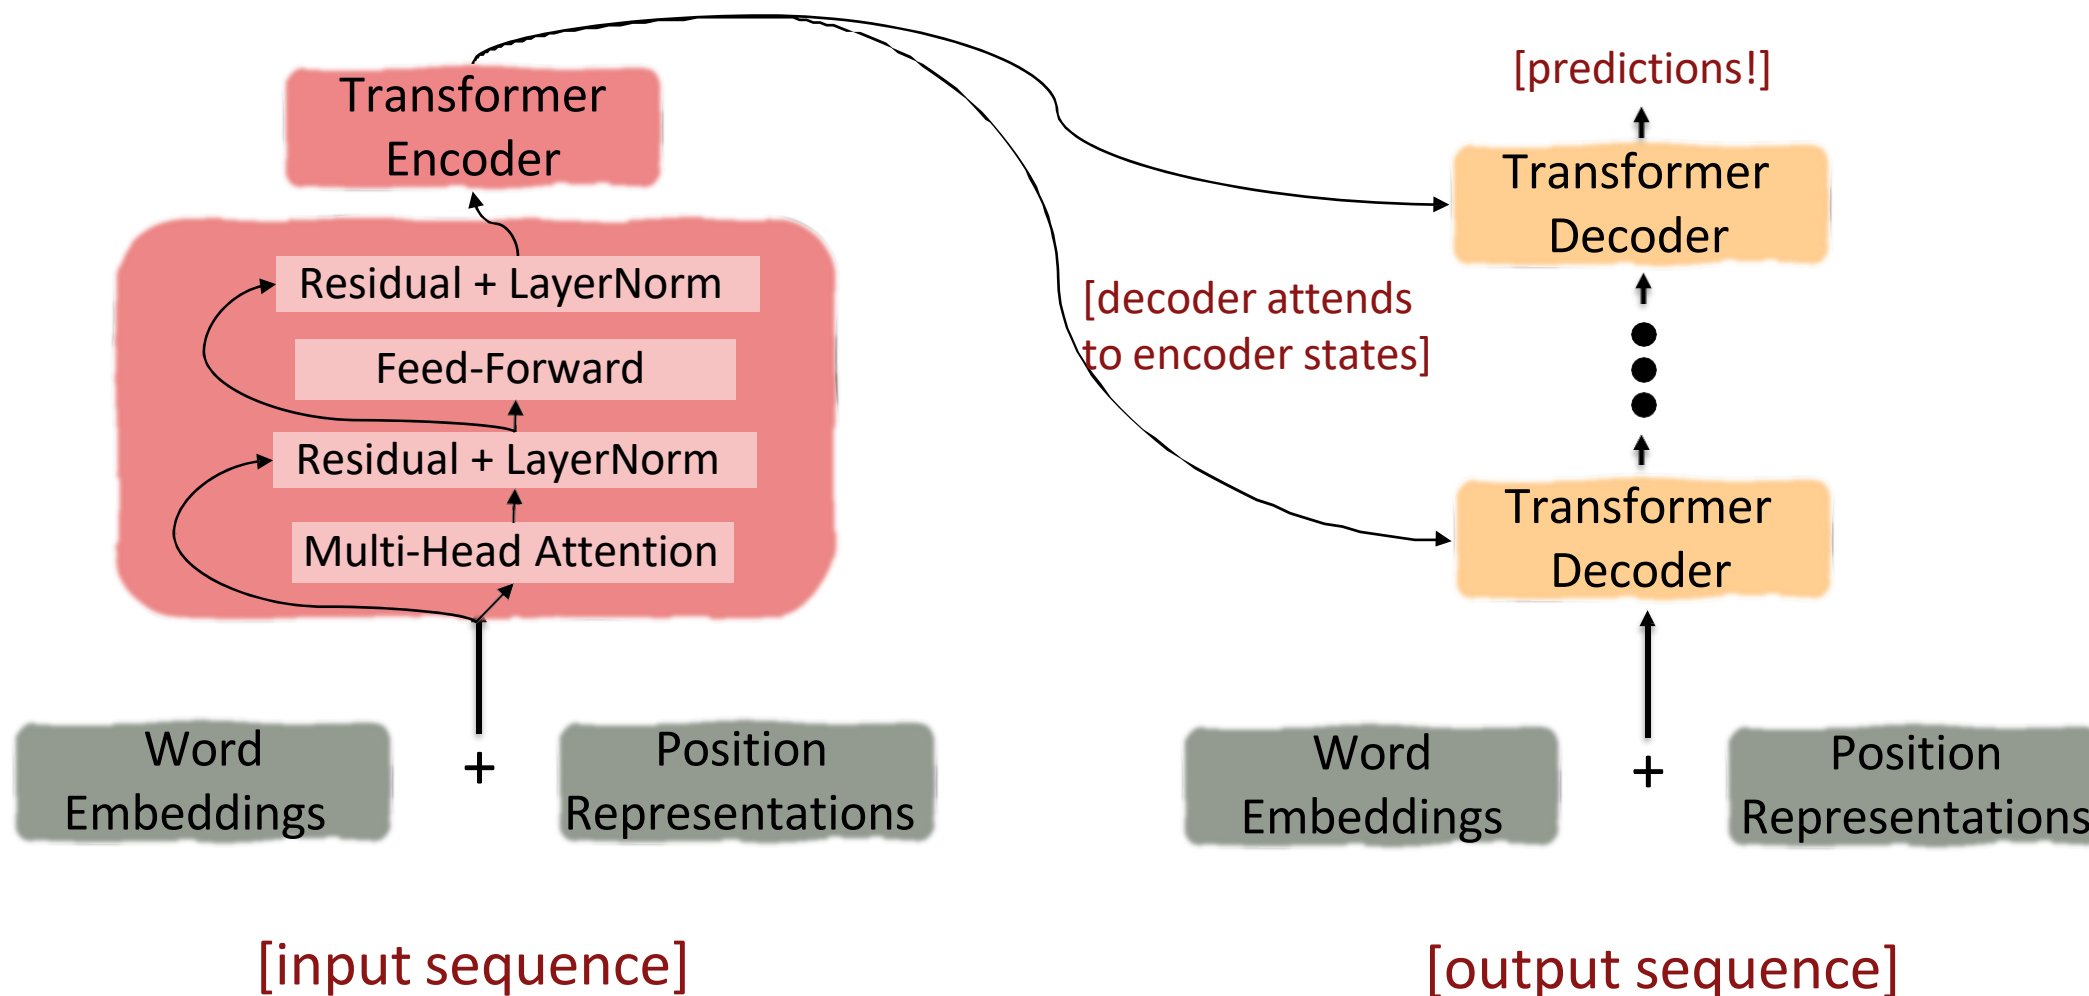
# The Transformer Encoder-Decoder [\[Vaswani et al., 2017\]](#)

Looking back at the whole model, zooming in on an Encoder block:



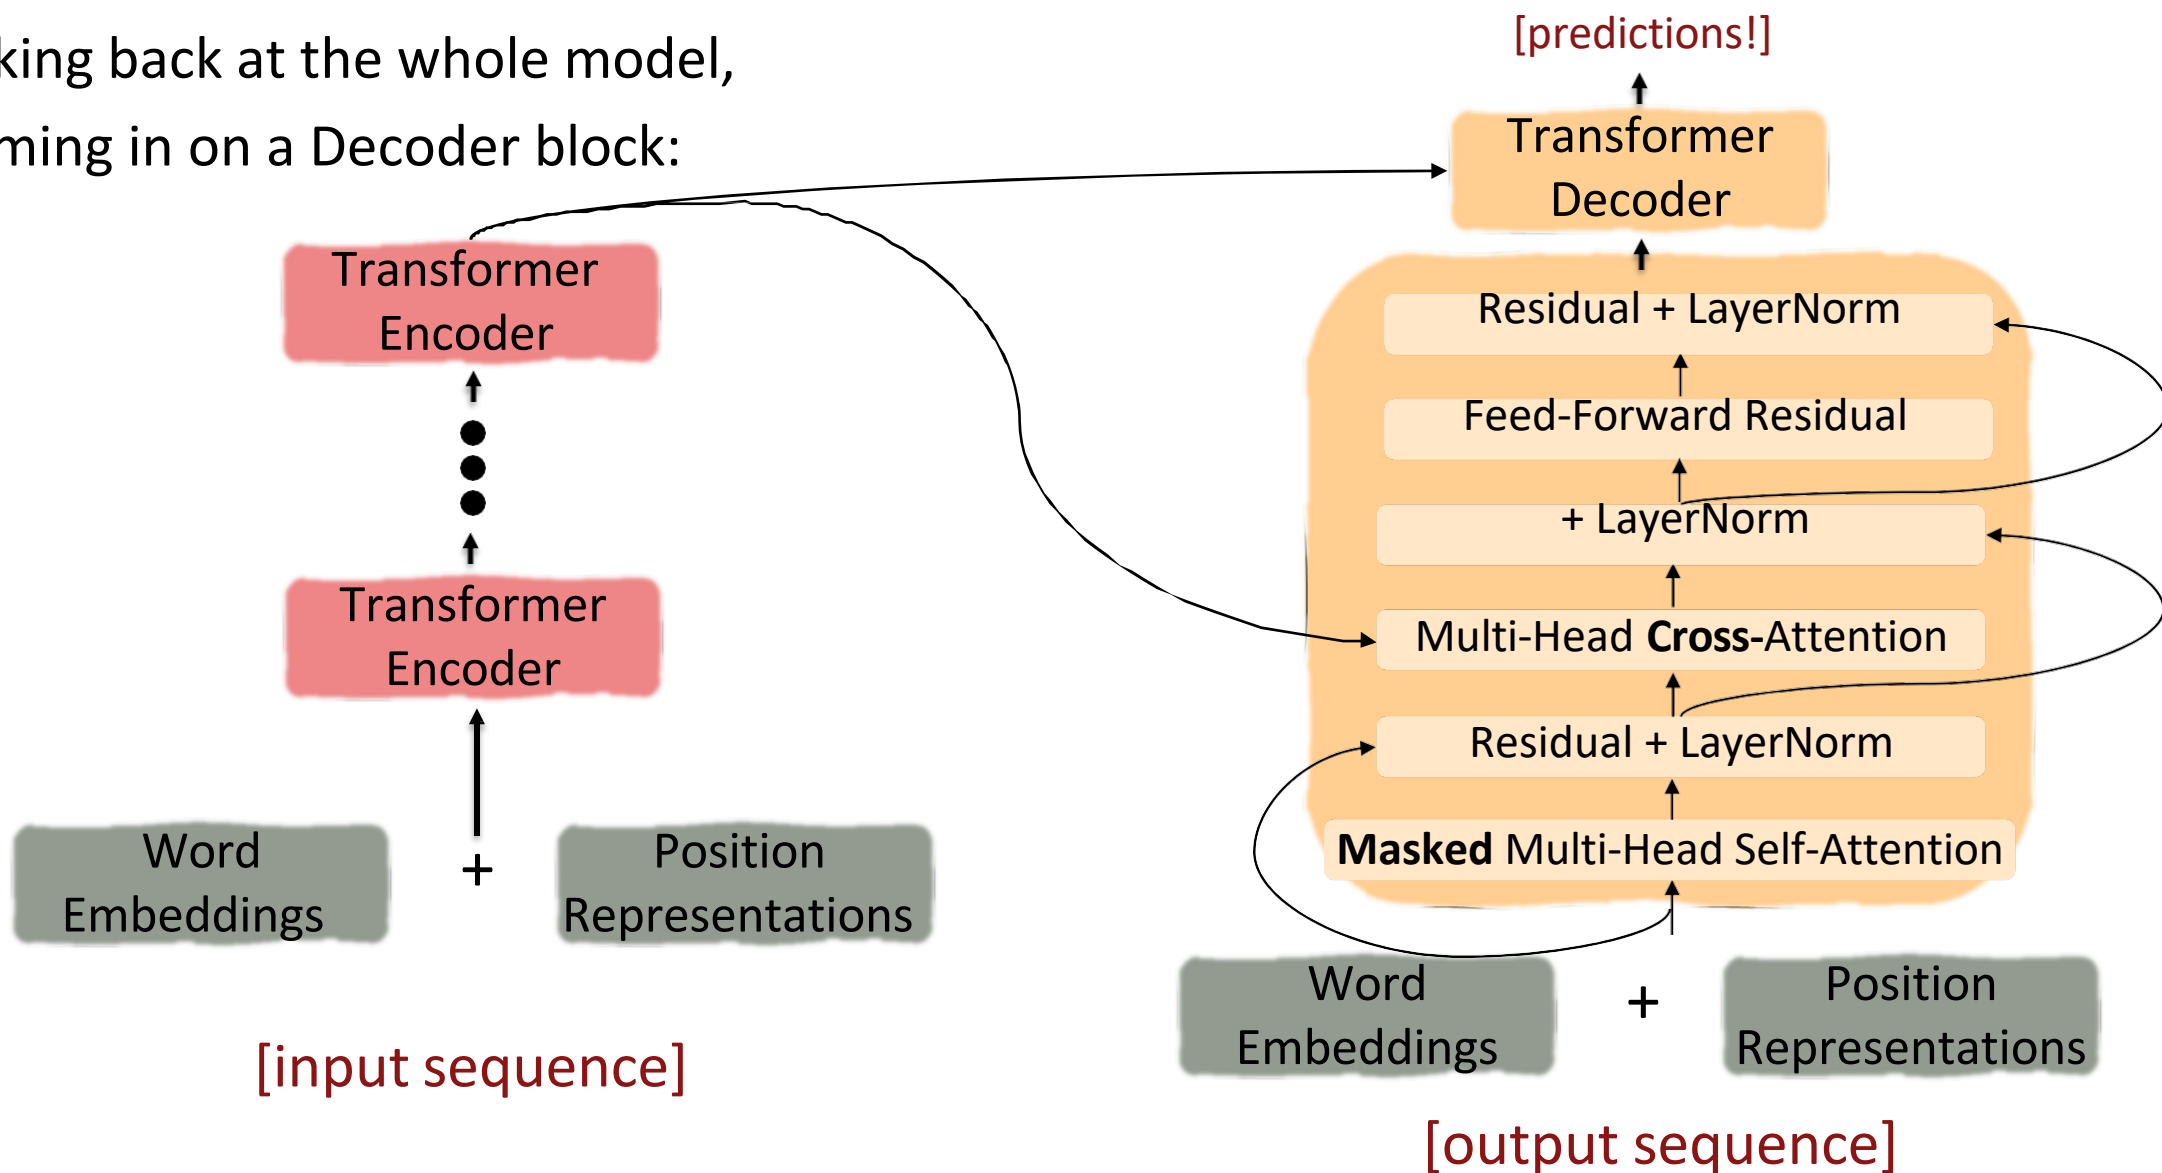
# The Transformer Encoder-Decoder [Vaswani et al., 2017]

Looking back at the whole model, zooming in on an Encoder block:



# The Transformer Encoder-Decoder [Vaswani et al., 2017]

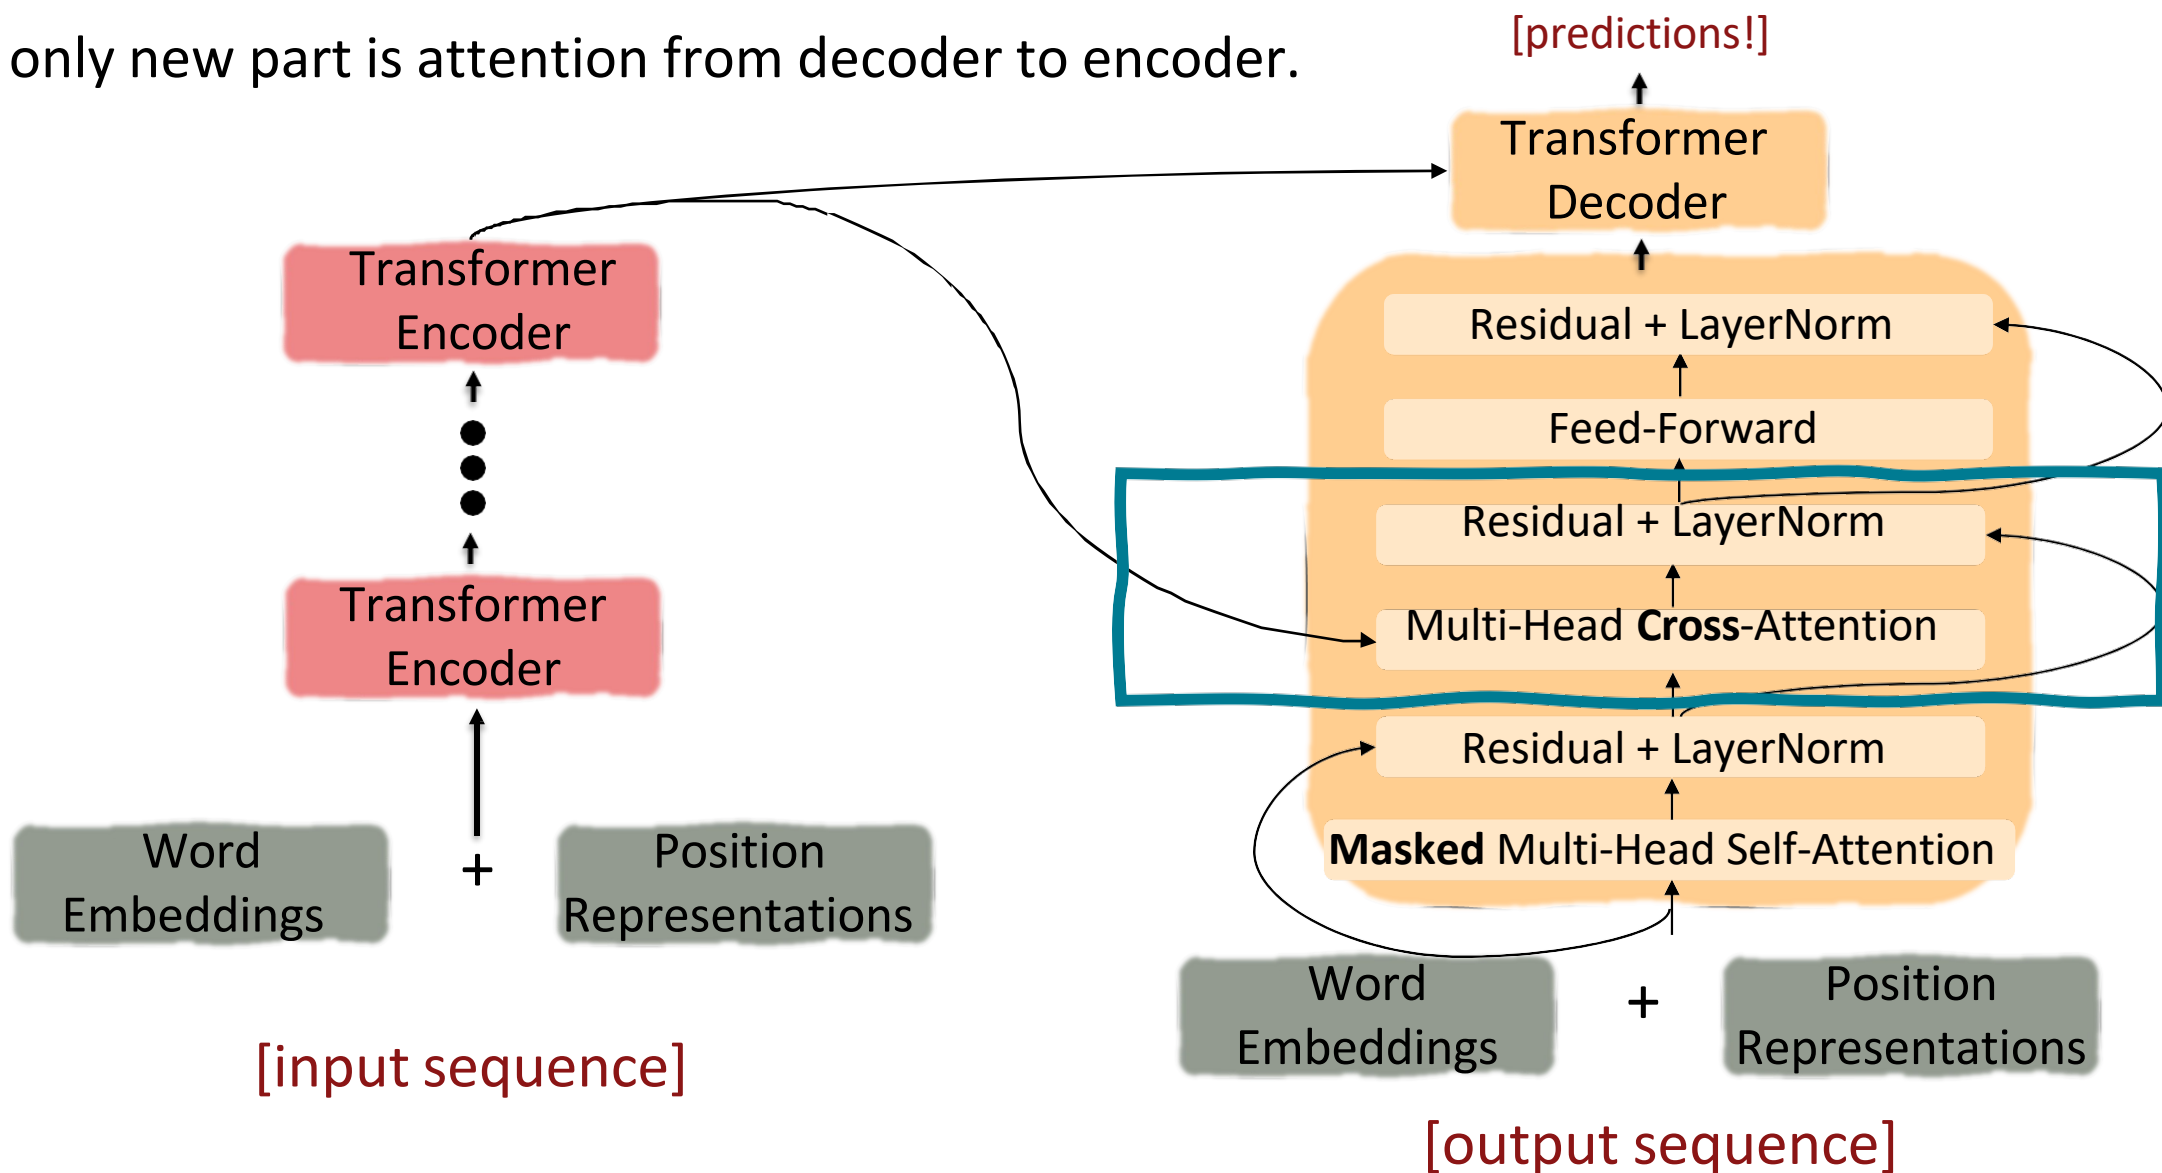
Looking back at the whole model,  
zooming in on a Decoder block:



# The Transformer Encoder-Decoder [Vaswani et al., 2017]

The only new part is attention from decoder to encoder.

Like



# The Transformer Decoder: **Cross-attention** (details)

- We saw that self-attention is when keys, queries, and values come from the same source.
- In the decoder, we have attention that looks more like what we saw last week.
- Let  $h_1, \dots, h_T$  be **output** vectors **from** the Transformer **encoder**;  $x_i \in \mathbb{R}^d$
- Let  $z_1, \dots, z_T$  be input vectors from the Transformer **decoder**,  $z_i \in \mathbb{R}^d$
- Then keys and values are drawn from the **encoder** (like a memory):
  - $k_i = Kh_i, v_i = Vh_i$ .
- And the queries are drawn from the **decoder**,  $q_i = Qz_i$ .

# The Transformer Encoder: Cross-attention (details)

- Let's look at how cross-attention is computed, in matrices.
  - Let  $H = [h_1; \dots; h_T] \in \mathbb{R}^{T \times d}$  be the concatenation of encoder vectors.
  - Let  $Z = [z_1; \dots; z_T] \in \mathbb{R}^{T \times d}$  be the concatenation of decoder vectors.
  - The output is defined as  $\text{output} = \text{softmax}(ZQ(HK)^\top) \times HV$ .

First, take the query-key dot products in one matrix multiplication:  $ZQ(HK)^\top$

$$ZQ \quad K^\top H^\top = ZQK^\top H^\top \in \mathbb{R}^{T \times T}$$

All pairs of attention scores!

Next, softmax, and compute the weighted average with another matrix multiplication.

$$\text{softmax} \left( ZQK^\top H^\top \right) HV = \text{output} \in \mathbb{R}^{T \times d}$$

# Outline

1. From recurrence (RNN) to attention-based NLP models
2. Introducing the Transformer model
3. Great results with Transformers

# Great Results with Transformers

First, Machine Translation from the original Transformers paper!

| Model                           | BLEU  |              | Training Cost (FLOPs) |                     |
|---------------------------------|-------|--------------|-----------------------|---------------------|
|                                 | EN-DE | EN-FR        | EN-DE                 | EN-FR               |
| ByteNet [18]                    | 23.75 |              |                       |                     |
| Deep-Att + PosUnk [39]          |       | 39.2         |                       | $1.0 \cdot 10^{20}$ |
| GNMT + RL [38]                  | 24.6  | 39.92        | $2.3 \cdot 10^{19}$   | $1.4 \cdot 10^{20}$ |
| ConvS2S [9]                     | 25.16 | 40.46        | $9.6 \cdot 10^{18}$   | $1.5 \cdot 10^{20}$ |
| MoE [32]                        | 26.03 | 40.56        | $2.0 \cdot 10^{19}$   | $1.2 \cdot 10^{20}$ |
| Deep-Att + PosUnk Ensemble [39] |       | 40.4         |                       | $8.0 \cdot 10^{20}$ |
| GNMT + RL Ensemble [38]         | 26.30 | 41.16        | $1.8 \cdot 10^{20}$   | $1.1 \cdot 10^{21}$ |
| ConvS2S Ensemble [9]            | 26.36 | <b>41.29</b> | $7.7 \cdot 10^{19}$   | $1.2 \cdot 10^{21}$ |



# Great Results with Transformers

Next, document generation!

| Model   | Test perplexity | ROUGE-L |
|---|-----------------|---------|
| <i>seq2seq-attention, <math>L = 500</math></i>                | 5.04952         | 12.7    |
| <i>Transformer-ED, <math>L = 500</math></i>                   | 2.46645         | 34.2    |
| <i>Transformer-D, <math>L = 4000</math></i>                   | 2.22216         | 33.6    |
| <i>Transformer-DMCA, no MoE-layer, <math>L = 11000</math></i> | 2.05159         | 36.2    |
| <i>Transformer-DMCA, MoE-128, <math>L = 11000</math></i>      | 1.92871         | 37.9    |
| <i>Transformer-DMCA, MoE-256, <math>L = 7500</math></i>       | 1.90325         | 38.8    |

The old standard



Transformers all the way down.



# Great Results with Transformers

Before too long, most Transformers results also included **pretraining**. Transformers' parallelizability allows for efficient pretraining, and have made them the de-facto standard.

On this popular aggregate benchmark, for example:



**All** top models are Transformer (and pretraining)-based.

| Rank Name |                          | Model                 | URL                 | Score |
|-----------|--------------------------|-----------------------|---------------------|-------|
| 1         | DeBERTa Team - Microsoft | DeBERTa / TuringNLRv4 |                     | 90.8  |
| 2         | HFL iFLYTEK              | MacALBERT + DKM       |                     | 90.7  |
| +         | 3                        | Alibaba DAMO NLP      | StructBERT + TAPT   | 90.6  |
| +         | 4                        | PING-AN Omni-Sinitic  | ALBERT + DAAF + NAS | 90.6  |
|           | 5                        | ERNIE Team - Baidu    | ERNIE               | 90.4  |
|           | 6                        | T5 Team - Google      | T5                  | 90.3  |

**More results Thursday when we discuss pretraining.**

[[Liu et al., 2018](#)]